

**OptiDBX Study Series · Document 5**

Page-by-Page Presentation Scripts

Exactly what to say and do on each page of the dashboard — Home, Scenarios, Live Session, Results & History, and Performance Study — with answers ready for every likely question.

Project: OptiDBX — Safe PostgreSQL Auto-Tuning

Document: 5 of the OptiDBX Study Series

Use with: the live dashboard open in front of you

Prepared: 22 September 2026

THE OPTIDBX STUDY SERIES — READ THESE TOGETHER**1 Demo Scenarios Study Guide**

Understand every scenario the dashboard can demonstrate.

2 Viva Question Bank

Questions a panel may ask, with grounded answers.

3 Technical Terms Glossary

Every term in the project, defined plainly.

5 Page-by-Page Presentation Scripts

What to say on each page of the dashboard.

YOU ARE HERE

BEFORE YOU START

Page-by-Page Presentation Scripts

A spoken walkthrough of the dashboard. Each page has four blocks: **Say this** (a script you can read aloud), **Do this** (clicks and where to point), **If asked** (quick answers with a pointer to the Viva Bank), and **Then** (how to move on). The orange trail shows where you are in the flow.

The one-line pitch (memorise this): “OptiDBX safely tunes one PostgreSQL knob on a live database — it changes something only when there's a real bottleneck, and keeps the change only when the measurements prove it helped.”

Setup checklist. Backend and dashboard running; on **Home** to start; optionally use *Clear simulated history* for a clean Results page; and, if you want to show a completed live run quickly, kick off a short Live Session a few minutes before you reach that page. Suggested full run: **7–9 minutes**.

Five pages, in order

Home → Scenarios → Live Session → Results & History → Performance Study.

1 Home

Set the story: what OptiDBX is and why it is trustworthy.

Home → Scenarios → Live Session → Results & History → Performance Study

Suggested timing: ~60–90 seconds

SAY THIS

This is OptiDBX — a safe auto-tuner for PostgreSQL. Its whole job is to watch a live database, make one small, reversible change when there is a genuine bottleneck, and keep that change only if the measurements prove it actually helped.

The headline is this number: it went from never keeping a beneficial change to keeping half of them — and every kept change is backed by a real, measured improvement. The tagline sums up the philosophy: Understand, Tune, Verify.

Below that we explain the signals it reads — the operating system's CPU, memory and disk, and the database itself — and the four steps it always follows: Monitor, Detect, Act safely, Verify.

DO THIS

- Start on the hero. Point at the **0 → 50%** statistic as you say the headline.
- Scroll slowly through **How it works** (four steps) and the three pillars — **Smart, Safe, Honest**.
- Do not rush; this page is the promise you will prove on the later pages.

IF ASKED

- “**What exactly does it tune?**” — One knob — `max_parallel_workers_per_gather`. Doc 2 Q2.1, Doc 3 term.
- “**Is 50% good?**” — Yes, and honest — it was 0% before. Full context in Doc 2 Q6.1–6.2.

THEN

Say: “Let me show you the situations it handles.” → click **Scenarios**.

2 Scenarios

Offer the menu of situations; pick one to walk through.

Home → **Scenarios** → Live Session → Results & History → Performance Study

Suggested timing: ~30 seconds on this page

SAY THIS

These nine cards are the situations OptiDBX can face. Four end in a change being kept, one is caught and reverted, two correctly do nothing, and the last two show it staying inside its safe scope and its worst-case recovery.

Each card shows its category, whether it is a simulated walkthrough or a live run, and the expected outcome. Let's open the clearest success case first.

DO THIS

- Point out the grouping: keeps, a rollback, the no-action cases, scope & recovery.
- Note the **Simulated** vs **Live** badges — call this out; the honesty matters.
- Click **CPU / Parallelism Contention** (Scenario 1) to begin.

IF ASKED

- **"Are the simulated numbers made up?"** — No — seeded from measured runs. Doc 1 honesty note, Doc 2 Q6.2.
- **"Which should I demo?"** — 1 (keep) → 5 (rollback) → 6 (live, no-action). Doc 1 Part C order.

THEN

A scenario opens in the runner. Narrate the pipeline as it plays, then **Download PDF report** at the end.

Home → **Scenarios** → Live Session → Results & History → Performance Study

Suggested timing: ~60 seconds per scenario

SAY THIS

Watch the pipeline across the top: Baseline, Detection, Apply, Observation, Decision, Report. On the left is the telemetry; on the right, a plain-language read-out of what the system is thinking at each step.

It takes a warm baseline, then needs three consecutive bad readings before it will act — no reacting to a single spike. It applies one step down, measures the owned workload, and only keeps the change because the tail latency genuinely fell. The outcome card shows the before-and-after, and every run produces a PDF report.

DO THIS

- Let the stepper advance; point at each stage as it lights up.
- At the end, read the **before/after bars** aloud (grey is before, orange is after).
- Click **Download PDF report** to show the artifact. Use **Replay** or go back for another scenario.

IF ASKED

- “**Why three readings?**” — Debounce transient spikes. Doc 1 Scenario 7, Doc 2 Q3.4.
- “**Why keep this one?**” — Net-benefit rule — measured tail-latency gain, no regression. Doc 2 Q4.5.

THEN

Say: “That was illustrative — now the same pipeline on a real database.” → click **Live Session**.

3 Live Session

Run the identical pipeline against the real PostgreSQL backend.

Home → Scenarios → **Live Session** → Results & History → Performance Study

Suggested timing: setup ~20 s, then the run (2–5 min)

SAY THIS

This is the real thing. I pick a pressure level — Low, Medium or High — and a duration, then start. From here the interface is deliberately identical to the scenario you just saw, but every number is now measured live on PostgreSQL.

While it runs, the controls lock so nothing disturbs the measurement, and I can pop open live metrics at any time. When the duration ends, it reaches a verdict on its own and produces the same before/after report — on a healthy database, the correct answer is often to change nothing, and it will say so.

DO THIS

- Choose a profile (Medium is a safe default) and a duration, then click **Start live session**.
- Point out the **session lock** and the **Live metrics** popup button.
- If presenting time is short, start a short run earlier and return to it here already finishing.

IF ASKED

- “**Why did it do nothing?**” — No sustained contention on a healthy host — correct. Doc 1 Scenario 6, Doc 2 Q6.4.
- “**Is it really live?**” — Yes — measured on the real backend; the badge says Live. Doc 2 Q8.5.

THEN

When it completes, say: “Every run is recorded.” → click **Results & History**.

4 Results & History

Show the audit trail and the exportable evidence.

Home → Scenarios → Live Session → **Results & History** → Performance Study

Suggested timing: ~30–45 seconds

SAY THIS

Everything that has run — live sessions and simulated walkthroughs — is listed here, clearly labelled, newest first. Open any row and you get its before-and-after evidence and a downloadable PDF report.

This is the accountability layer: nothing is hidden, live and simulated are never mixed up, and every claim on the earlier pages can be traced back to a specific run right here.

DO THIS

- Point at the **Live** vs **Simulated** badges in the table.
- Open the live session you just ran; read its before/after and click **Download PDF**.
- Mention **Clear simulated history** exists for a clean start before a demo.

IF ASKED

- “Where are live results stored?” — In PostgreSQL; simulated ones are local and labelled. Doc 2 Q8.5.
- “Can I trust the keeps?” — Yes — the same log shows the rollbacks and no-actions. Doc 2 Q9.3.

THEN

Say: “If you want a rigorous head-to-head, there’s a dedicated study.” → click **Performance Study**.

5 Performance Study

Offer the fair, paired baseline-vs-tuned comparison for the sceptic.

Home → Scenarios → Live Session → Results & History → **Performance Study**

Suggested timing: ~45 seconds to explain (a full study runs for minutes)

SAY THIS

This page is for the sceptic. It runs a fair, paired comparison: the same owned workload with unchanged settings versus with OptiDBX tuning, repeated and interleaved to cancel out ordering and cache effects.

You fix the success criteria before the run — so the result can't be rationalised afterwards — and it reports queries per second, median and p95 latency for each side, with a clear verdict. Every claim stays linked to the exact runs that produced it.

DO THIS

- Point at the two presets (**Pilot** for a quick look, **Evidence study** for the rigorous one).
- Show the config: workload, paired repetitions, warm-up, measurement seconds, starting parallelism.
- Note **success criteria fixed before the run** and that safety evidence is kept separate from speed.
- For a live demo use the Pilot preset; a full Evidence study is best run beforehand.

IF ASKED

- **"Why paired and interleaved?"** — To cancel order/cache bias. Doc 2 Q4.3, Doc 3 'warm-up exclusion'.
- **"Isn't the regime constructed?"** — Yes, and we say so — Doc 2 Q6.2, report §8.

THEN

Close here, or loop back to Home for the summary line.

Closing & handling questions

Closing line. “So: it understands the system, tunes one thing safely, and verifies the result — keeping the change only when the numbers earn it, and doing nothing when nothing is wrong. That honesty is the point.”

Three reflexes for tough questions:

- **Volunteer the caveat first.** The 50% keep rate is real but from a constructed, small-sample regime — say so before you're asked (Doc 2 Q6.2).
- **Point to the evidence.** Any claim can be opened in Results & History or reproduced in the Performance Study — don't argue, demonstrate.
- **Fall back to the bank.** Every likely question has a grounded answer in Document 2; every term is in Document 3.

OptiDBX Study Series · Document 5 — Page-by-Page Presentation Scripts. Scripts match the current dashboard; pointers reference the companion documents. Companion documents: 1 Demo Scenarios Study Guide, 2 Viva Question Bank, 3 Technical Terms Glossary.